

Lecture 5: Fingerprinting

David Woodruff


How to Pick a Random Prime

- How to pick a random prime in the range $\{0, 1, \dots, M-1\}$?
 - Pick a random integer X in the range $\{0, 1, \dots, M-1\}$
 - Check if X is a prime. If so, output it. Else go back to the first step
- How to check if X is prime?
 - Miller-Rabin primality test very efficient but fails with tiny probability
 - Agrawal-Kayal-Saxena has a worse running time, but deterministic
- How likely is X to be prime?

Density of Primes

- Let $\pi(n)$ be the number of primes in the set $\{1, 2, \dots, n\}$
- **Prime Number Theorem:** $\lim_{n \rightarrow \infty} \frac{\pi(n)}{n / \ln n} = 1$
- Chebyshev: $\pi(n) > n / \ln n$ for every $n \geq 2$
 - If we want at least k primes in $\{1, 2, \dots, n\}$, then $n \geq 2k \lg k$, if $k \geq 4$
- Dusart: For $n > 60184$, we have $\frac{n}{\ln n - 1.1} > \pi(n) > \frac{n}{\ln n - 1}$

String Equality Problem



- x and y are N -bit strings
- Alice and Bob want to exchange messages to decide if $x = y$
- Alice could send x to Bob but this takes N communication
 - Is there a more efficient scheme?

String Equality Problem

- Suppose we are OK if we achieve a probabilistic guarantee:
 - If $x = y$, then $\Pr[\text{Bob says } \mathbf{equal}] = 1$
 - If $x \neq y$, then $\Pr[\text{Bob says } \mathbf{unequal}] \geq 1 - \delta$
- Protocol
 - Alice chooses a random prime p from $\{1, \dots, M\}$ for $M = \lceil 2 \cdot (5N) \cdot \lg(5N) \rceil$
 - She sends Bob p and the value $h_p(x) = x \bmod p$, where we think of x as an integer in $\{0, 1, 2, \dots, 2^N - 1\}$
 - If $h_p(x) = y \bmod p$, Bob says **equal**, else he says **unequal**

String Equality Problem

- **Lemma:** If $x = y$, then Bob always says **equal**
- **Proof:** If $x = y$, then $x \bmod p = y \bmod p$. So Bob's test will always succeed
- **Lemma:** If $x \neq y$, then $\Pr[\text{Bob says } \mathbf{equal}] \leq .2$
- **Proof:** Interpret $x, y \in \{0, 1, 2, \dots, 2^N - 1\}$

If Bob says **equal**, then $x \bmod p = y \bmod p$, i.e., $(x-y) = 0 \bmod p$

So p divides $D = |x-y|$, and $D < 2^N$

$D = p_1 \cdot p_2 \cdots p_k$ for primes p_1, \dots, p_k which may repeat

Since each $p_i \geq 2$, we have $k < N$

$$\Pr[p \text{ divides } D] \leq \frac{N}{\text{number of primes in } \{1, 2, \dots, M\}} \leq \frac{N}{5N} = \frac{1}{5} \quad \text{why?}$$

Communication Cost

- If Alice were to naively send x to Bob, would take N bits of communication
- Instead she sends a prime p and $x \bmod p$, where p is in $\{1, 2, \dots, M\}$ and $M = \lceil 2 \cdot (5N) \cdot \lg(5N) \rceil$
- Communication = $O(\log p) = O(\log M) = O(\log N + \log \log N) = O(\log N)$ bits

Reducing the Error Probability

- We have 20% error probability, how to reduce it to δ ?
- Repeat the scheme $r = \log_5(\delta^{-1})$ times independently with primes $p_1, \dots, p_r \in \{1, 2, \dots, M\}$, and $M = \lceil 2 \cdot (5N) \cdot \lg(5N) \rceil$
 - Bob outputs **equal** if and only if $x = y \bmod p_i$ for each i
 - If $x = y$, Bob outputs **equal** with probability 1
 - If $x \neq y$, Bob outputs **equal** with probability at most $\left(\frac{1}{5}\right)^{\lg_5(\frac{1}{\delta})} \leq \delta$
 - Communication cost is $O(\log(1/\delta) \log N)$. **Can we do better?**
- If instead Alice sets $M = 2 \cdot sN \lg(sN)$, the number of primes in $\{1, 2, \dots, M\}$ is at least sN , and so error probability is $1/s$. Set $s = 1/\delta$.
 - Communication is $O(\log M) = O(\log s + \log N) = O(\log(1/\delta) + \log N)$

Fingerprinting (the Karp-Rabin Method)

- In the string-matching problem, we have
 - A text T of length m
 - A pattern P of length n
- **Goal:** output all occurrences of the pattern P inside the text T
 - If $T = \text{abracadabra}$ and $P = \text{ab}$, the output should be $\{0,7\}$

abracadabra

- Consider $h_p(x) = x \bmod p$ for $x \in \{0,1\}^n$, where we think of x as an integer in $\{0, 1, 2, \dots, 2^n-1\}$

Fingerprinting (the Karp-Rabin Method)

- $h_p(x) = x \bmod p$ for $x \in \{0,1\}^n$
- Create x' by dropping the most significant bit of x , and appending a bit to the right
 - E.g., if $x = 0011001$, then x' could be 0110010 or 0110011
- Given $h_p(x) = z$, can we compute $h_p(x')$ quickly?
- Suppose x'_{lb} is the lowest-order bit of x' , and x_{hb} is the highest order bit of x
- $x' = 2(x - x_{hb} \cdot 2^{n-1}) + x'_{lb}$
- Since $h_p(a + b) = (h_p(a) + h_p(b)) \bmod p$, and $h_p(2a) = 2h_p(a) \bmod p$,
$$h_p(x') = (2h_p(x) - x_{hb} \cdot h_p(2^n) + x'_{lb}) \bmod p$$
- *Given $h_p(x)$ and $h_p(2^n)$, this is just $O(1)$ arithmetic operations mod p*

Fingerprinting (the Karp-Rabin Method)

- $T_{a...b}$ denotes the string from the a -th to b -th positions of T , inclusive
 - **Goal:** output all locations a in $\{0, 1, \dots, m-n\}$ such that $T_{a...a+(n-1)} = P$
1. Pick a random prime $p \in \{1, 2, \dots, M\}$ with $M = \lceil 2s n \lg(sn) \rceil$ for some s
 2. Compute $h_p(P)$ and $h_p(2^n)$ and store the results
 3. Compute $h_p(T_{0...n-1})$ and check if it equals $h_p(P)$. If so, output **match** at location 0
 4. For each $i \in \{0, \dots, m-n-1\}$, compute $h_p(T_{i+1...i+n})$ using $h_p(T_{i...i+n-1})$ and $h_p(2^n)$. If $h_p(T_{i+1...i+n}) = h_p(P)$, output **match** at location $i+1$

Error Probability

- $m - n + 1 \leq m$ comparisons, each with probability at most $1/s$ of failure
- By a union bound, the probability there is at least one failure is at most m/s
- If $s = 100m$, we succeed on all comparisons with probability $\geq 99/100$
- $M = \lceil 2s n \lg(sn) \rceil = O(mn \log(mn))$, so $O(\log m + \log n)$ bits to store
- Since p in $\{1, 2, \dots, M\}$, p takes $O(\log m + \log n)$ bits to store
- Assume unit-cost RAM model, so operations on $O(\log(mn))$ bits take $O(1)$ time

Running Time

- Computing $h_p(x)$ for n -bit x can be done in $O(n)$ time. **Why?**
 - Generate hashes of powers of 2 by multiplying each previous hash by 2 mod p
 - Then can use shifting
- So $h_p(P)$, $h_p(2^n)$, and $h_p(T_{0,\dots,n-1})$ can be computed in $O(n)$ time
- Computing $h_p(T_{i+1\dots i+n})$ using $h_p(T_{i\dots i+n-1})$ and $h_p(2^n)$ can be done in $O(1)$ time!
- Total time is $O(m + n)$, which is optimal

Fingerprinting Extensions

- Fingerprinting also works for strings $x \in \{0, 1, 2, \dots, q - 1\}^n$
- Think of x as an integer $\sum_{i=0, \dots, n-1} q^i \cdot x_i$ in its q -ary representation
- Drop the leftmost digit of x to create x' , and append a digit to the right
 - If $x = x_{n-1}, x_{n-2}, x_{n-3}, \dots, x_0$, then $x' = x_{n-2}, x_{n-3}, \dots, x_0, x'_0$
- $x' = q(x - x_{n-1} \cdot q^{n-1}) + x'_0$
- $h_p(x') = (q \cdot h_p(x) - x_{n-1} \cdot h_p(q^n) + x'_0) \bmod p$
- Given $h_p(x)$ and $h_p(q^n)$, computing $h_p(x')$ requires $O(1)$ arithmetic operations mod p (assuming $q \bmod p$ can be computed quickly)

Extensions

- How would you solve the following?
- Given an $m_1 \times m_2$ -bit rectangular binary text T , and an $n_1 \times n_2$ -bit pattern P , where $n_1 \leq m_1$ and $n_2 \leq m_2$, find all occurrences of P inside T . Show how to do this in $O(m_1 m_2)$ time
- Assume you can do modular arithmetic of integers at most $\text{poly}(m_1 m_2)$ in $O(1)$ time

Extensions

- Walk through the columns of T , and create fingerprints $h_q(T_{[i,i+n_1-1],j})$ of the n_1 values

$$T_{i,j}, T_{i+1,j}, \dots, T_{i+n_1-1,j}$$

- $q \leq \text{poly}(m_1 m_2 n_1)$
- Walk through the rows of T , and for the (i,j) -th entry, create a fingerprint of the n_2 values

$$h_q(T_{[i,i+n_1-1],j}), h_q(T_{[i,i+n_1-1],j+1}), \dots, h_q(T_{[i,i+n_1-1],j+n_2-1})$$

- **Note:** the fingerprints are of q -ary instead of binary strings, but when fingerprinting these strings we can use a prime $p \leq \text{poly}(m_1 m_2 n_1 n_2)$. **Show this!**
- Walking through the columns and rows and creating the fingerprints, and comparing with the hash of the pattern P , takes $O(m_1 m_2)$ time